

WebRetriever — 正式比赛 FAQ

English version below

一、赛程与提交流程

时间节点

Q: 8月27日之前需要做什么？

8月27日（正式评测开启）前，每支队伍需通过一次冒烟测试确保环境已安装和流程通过。冒烟测试通过仅表示你的输出符合比赛格式要求。正式提交评测之前，欢迎继续调优你的 Agent。

Q: 正式提交窗口期是具体什么时间？以什么方式算提交时间？

正式提交评测的开启时间为8月27日10:00（北京时间）- 9月2日23:59（北京时间）。以选手在专属提交群中发起评测（而非冒烟测试）的时间为准。

Q: 备赛期间与正式提交窗口的提交有什么区别？

备赛期间提交代码至私有仓库，可以通过专属提交群内的EvalBot进行冒烟测试，验证输出格式是否符合比赛要求。

正式提交窗口期在冒烟测试验证合格后可以发起比赛评测对每一题进行评分。比赛评测仅在正式提交窗口期间开放，选手在备赛期可以用开源数据集

🔗 <https://huggingface.co/datasets/Mininglamp-2718/WebRetriever> 的协议 3 题目进行 Agent 调优。

冒烟测试

Q: 冒烟测试会消耗正式评测次数吗？

不会。冒烟属于「提交代码」的验证流程，与正式评测是两条独立的路，没有每日限制。每日评测次数只会被「正式评测（跑分）」消耗。

Q: 冒烟测试是只做语法/依赖安装，还是会真实跑任务？

冒烟测试会真实跑少量任务（非比赛正式任务），确保参赛者代码能够按照规定进行输入输出，但是不会对任务执行结果进行评分。

Q: 冒烟测试通过是看最新一次，还是只要成功过就行？

每次代码改动后，评测前都需要重新通过冒烟测试。不是“曾经通过就行”。

Q: 怎么算冒烟测试通过？

在你的专属提交群内，@WR-EvalBot提交代码后。收到Bot发送的"冒烟测试通过，环境就绪！可以发起评测了"或"代码验证通过，环境就绪！可以发起评测了"即代表冒烟测试通过。

或直接询问WR-EvalBot，"冒烟测试通过了吗？"即可获得冒烟测试结果。

正式评测

Q: 每日正式评测次数是几次？

1 次。评测次数刷新时间为00:00（北京时间），以在专属提交群中发起评测（而非冒烟测试）的时间为准。

Q: 评测跑的是全量题还是随机子集？

每次评测跑的都是全量题，总共 100 道。

Q: 所有参赛队伍用同一套试题吗？

是的，比赛期间所有队伍使用同一套试题。

Q: 每次提交评测后是否会有本次评测分数反馈？

会，每次评测后EvalBot会经过评估后反馈分数。请注意，在赛中bot自动反馈的分数仅为任务答案准确得分，即任务答案是否正确的分数，不包括轨迹评分。在评测窗口结束，全部成绩复核后结合轨迹正确性才会公布最终成绩分数。

Q: 正式提交期间能否看到其他队伍的分数/排名？

比赛期间会每日中午12:00在正式比赛队长群中匿名公布当前累积排名第一的分数。

同上，赛中排名公布的分数仅为答案正确的得分，最终分数会结合轨迹复核后更新，最终分数为轨迹正确+答案正确的修正得分。

二、沙箱环境与基础设施

浏览器环境

Q: 浏览器沙箱环境是什么？屏幕分辨率多少？

浏览器沙箱环境为腾讯云沙箱，屏幕分辨率默认 1920 × 1080。

Q: 每道题是否从干净浏览器状态开始？（无历史 cookie/localStorage）

是的。

Q: 能否开多个 tab/新窗口？弹窗（window.open）是否被拦截？

可以开多个 tab。弹窗不会拦截，这属于真实网络环境的一部分，需要 Agent 自行解决。组委会会尽量保证出题网站是弹窗少的。

运行环境

Q: 可以在沙箱内使用 bash 工具吗？

不可以。评测环境是浏览器沙箱，无法使用 bash 工具（如 curl 等命令行操作）。

Q: 可以在 conda 环境中安装 JDK 吗？

可以。代码运行在独立的conda环境中（可自定义依赖），浏览器交互通过CDP连接组委会提供的云端沙箱。建议先在自己机器上验证自定义环境在 conda 下是否能稳定运行。

Q: 编程语言或开源框架有限制吗？

没有限制，但是选手需要按照比赛模板要求进行输入输出接口标准化，并且代码需要能够在conda环境中稳定运行。

Q: 沙箱内下载文件存哪？临时目录路径？磁盘空间限制？

下载操作需要调用 playwright 支持的下载方式存到代码目录下，路径由选手自定义。磁盘空间没有限制。

网络与安全

Q: 外网网站加载不出来，比赛中会出现吗？

比赛期间的评测沙盒环境由组委会提供，会尽量避免网站被墙/地区限制的问题。

Q: 我的 API 服务只有 IP 没有域名，域名白名单怎么办？

如果你的代理最终访问的是 OpenAI 官方 API 服务，组委会已对 OpenAI 的海外访问进行加速。可以在 config.json 中直接配置调用 OpenAI 官方 API 服务，无需额外购买域名。

请提供您的 Agent 运行时实际请求的算法服务入口域名，用于加入网络白名单。提交表单：
<https://ml.jsjform.com/f/jXbK5J>

Q: API Key 需要明文写进 Git 仓库吗？代码在仓库里安全吗？

私有仓库内容只有你的队伍可以看到，评测环境有严格的安全策略，可以保障 API Key 安全。

赛事结束后，依照选手个人意愿，主办方会将私有仓库进行销毁或转接给选手本人。

三、比赛规则与限制

步数与执行

Q: 单题有时间限制吗？

评测时不规定单题的时间限制，选手需根据评测总时长合理规划。

Q: max steps 具体是多少步？

官方规定最多100步，选手可以结合评测总时长限制和 Agent 策略自定义最大步长，不超过100步即可。

Q: 如何定义“执行一步”？截图+思考+动作=1步？还是每次 action 调用=1步？

Web Agent 基于当前网页内容预测动作并执行操作，就算一步。

Q: 超过 max steps 后是直接判失败还是取当前状态评分？

超过100步直接判失败。

Q: 有单次模型调用的超时限制、调用次数限制吗？

单次模型调用不超过 3 分钟，调用次数不限制，模型需要支持 8 进程并发请求调用。

Q: 有评测整体超时限制吗？

评测总时长不超过 4 小时。

操作边界

Q: 可以导航到搜索引擎进行搜索吗？

不可以使用外部搜索引擎。可以使用网站内搜索。

Q: 直接 CDP 操作是否合规？

直接 cdp 操作不合规。需要基于 playwright 创建 cdp session，基于此再执行 cdp 操作。

Q: 涉及 PDF 等文件的任务，可以下载后用 coding agent 处理吗？

可以。导航到页面后，下载对应数据并分析得到答案，整体算一步。

Q: 比赛会出现需要登录才能完成的任务吗？

不会。赛事题目会尽量避免需要登录的情况。开源数据集中如出现需要登录的情况，可能是因为网页发生了变化。

Q: 比赛会出现人机验证（验证码）吗？

反爬/风控验证（如点击验证）：可能出现，这属于真实网络环境的一部分，需要 Agent 自行解决。组委会会尽量保证出题网站风控级别低。

模型与合规

Q: 参赛队伍用的模型有什么限制？

Agent 必须通过 OpenAI 兼容格式的 API 调用模型。两种选择：

选项 A：商业闭源 API

各厂商允许使用的最高版本：

- OpenAI — GPT-5.4
- Anthropic — Claude 4.6
- Google — Gemini 3.1
- xAI — Grok 4.3
- 智谱 AI — GLM-5V-Turbo
- Moonshot — Kimi-K2.6
- 阿里云 — Qwen3.7 (新增)

超出上述版本的禁止使用。未在上表列出的厂商，只允许使用 2026 年 7 月 16 日之前发布的模型（以厂商官方公开发布公告日期为准），之后发布的禁止使用。

选项 B：自研/自部署模型

允许，赛后需提交模型权重供组委会校验。

组委会会审查模型调用情况，严查套壳和中转绕版本限制，违规成绩无效。

Q: 多个模型混用（比如 GPT 做推理、Claude 做 OCR）是否允许？

允许，但是模型都要符合组委会的版本限制。

Q: 模型 API 调用如果遇到 429 限流或超时，这步算失败还是可以重试？

可以重试，这个需要选手自己在agent设计中实现。

Q: 一个队伍可以提交多个不同 Agent 版本来 A/B 测试吗？

不限制提交代码的版本和次数，但是每日提交评测次数仅限1次。

四、评分标准

评分与排名逻辑

Q: 自动评测分数与最终公示分数有什么区别？

自动评测分数：在赛中bot自动反馈的分数仅为任务答案准确得分，即任务答案是否正确的分数，不包括轨迹评分。

最终公示分数：在评测窗口结束，全部成绩复核后结合轨迹正确性才会公布最终成绩分数。

Q: 正式比赛的最终公示分数评分标准是什么？

比赛采用 Protocol III（导航 + 取数）评分，判断标准为：

1. 是否导航到目标页面

2. 是否取出正确的字段（语义一致即可）
最终会有人工审核判断是否符合要求。

Q: 最终公示排行榜排名的依据是什么？是取最高分还是最后一次分数？

最终公示排名按照综合评分排序，综合评分为任务答案通过且轨迹评判通过的综合得分，分高者优先排名；同一评分下按总任务步数少者优先排；同一评分同一步数下按总时长少者优先排。

每支队伍如有多次提交，取最高综评分进入排名。

Q: 所有队伍的所有成绩都会复议吗？

是的。主办方会复查每支队伍的成绩，会结合轨迹+取数评分后在复议期结束后反馈给每支队伍最高评分的成绩。

Q: 判题用的是参赛者自己的模型还是统一模型？

参赛者只负责设计自己的 Agent，评分流程由组委会统一执行，判题模型不需要参赛者配置。

答案要求

Q: agent_answer 的格式要求是什么？纯文本？JSON？多字段怎么填？

纯文本。多字段可以\n或者空格进行间隔。

Q: "语义一致"的松紧度——数字要完全匹配还是允许四舍五入？中英文/简繁体是否等价？

以题目本身要求和目标页面的数值为准。等价，对语言没有限制。

Q: 部分完成有分吗？（导航对了但取数错，vs 导航错了但蒙对答案）

只有导航和取数都正确才算任务完成，其余情况都算失败。

五、常见问题排查

冒烟测试失败原因

#	原因	解决方案
1	提供 CDP_URL 后无法连接沙箱环境	保留比赛模板仓库中连接浏览器的兼容逻辑
2	个别任务的 capture.json 未正常捕获网络请求	检查网络请求捕获逻辑
3	连接沙箱后清空了标签页导致浏览器关闭	不要清空所有标签页
4	模型调用失败	检查 API Key 配置和网络连通性
5	执行过程中没有继续响应导致的超时	检查模型调用是否有超时重试机制

本地与沙箱差异

Q: 本地 Playwright/Puppeteer 跑得通，到沙箱行为不一致怎么办？

沙箱连接相比本地连接有些许差异，比赛模板仓库的初始化浏览器和连接浏览器部分已进行兼容，务必保留该逻辑。

WebRetriever — Competition & Evaluation FAQ

1. Schedule & Submission Process

Timeline

Q: What do I need to do before August 27?

Before August 27 (when formal evaluation opens), each team must pass smoke testing at least once to confirm that the environment is set up and the workflow is functional. Passing smoke testing only means your output meets the competition format requirements. You're welcome to continue tuning your Agent before the formal submission window opens.

Q: What is the exact formal submission window? How is submission time determined?

The formal evaluation window is August 27, 10:00 AM — September 2, 11:59 PM (Beijing Time). Submission time is based on when you initiate an evaluation (not smoke testing) in your dedicated submission group.

Q: What's the difference between submissions during the preparation period and the formal submission window?

During the preparation period, you can push code to your private repo and run smoke tests via EvalBot in your dedicated submission group to verify output format compliance. During the formal submission window, after passing smoke testing you can initiate competition evaluation to score each question. Competition evaluation is only available during the formal window. During preparation, you can use the open-source dataset at <https://huggingface.co/datasets/Mininglamp-2718/WebRetriever> (Protocol III questions) to tune your Agent.

Smoke Testing

Q: Does smoke testing consume formal evaluation attempts?

No. Smoke testing is part of the code submission verification process — it's a separate track from formal evaluation with no daily limit. Only formal evaluations (scoring runs) consume your daily quota.

Q: Does smoke testing just check syntax/dependencies, or does it actually run tasks?

Smoke testing runs a small number of real tasks (not formal competition tasks) to ensure your code handles input/output correctly, but it does not score task execution results.

Q: Is smoke test pass based on the latest run, or is passing once enough?

You need to pass smoke testing again after every code change before evaluation. It's not "once passed, always valid."

Q: How do I know if smoke testing passed?

After submitting code by @WR-EvalBot in your dedicated submission group, if you receive "Smoke test passed, environment ready! You can now initiate evaluation" or "Code verification passed, environment ready! You can now initiate evaluation," your smoke test has passed.

You can also ask WR-EvalBot directly: "Did smoke testing pass?"

Formal Evaluation

Q: How many formal evaluation attempts per day?

1 attempt. The daily quota resets at 00:00 (Beijing Time). The time is based on when you initiate evaluation (not smoke testing) in your dedicated submission group.

Q: Does evaluation run the full question set or a random subset?

Every evaluation runs the full set — 100 questions total.

Q: Do all teams get the same set of questions?

Yes, all teams use the same question set during the competition.

Q: Will I receive score feedback after each evaluation submission?

Yes. After each evaluation, EvalBot will provide a score. Note: the score provided by the bot during the competition is only the task answer accuracy score (whether answers are correct), and does not include trajectory scoring. After the evaluation window closes, final scores combining trajectory correctness will be published after review.

Q: Can I see other teams' scores/rankings during the competition?

During the competition, the cumulative top-ranked score will be published anonymously at 12:00 noon daily in the official team leader group.

Note: scores published during the competition are answer accuracy scores only. Final scores will be updated after trajectory review — the final score combines trajectory correctness + answer correctness.

2. Sandbox Environment & Infrastructure

Browser Environment

Q: What is the browser sandbox? What's the screen resolution?

The browser sandbox runs on Tencent Cloud. Default screen resolution is 1920 × 1080.

Q: Does each question start with a clean browser state? (No cookies/localStorage)

Yes.

Q: Can I open multiple tabs/new windows? Are popups (window.open) blocked?

Yes, you can open multiple tabs. Popups are not blocked — this is part of the real web environment and your Agent needs to handle it. The organizers will try to select websites with fewer popups.

Runtime Environment

Q: Can I use bash tools inside the sandbox?

No. The evaluation environment is a browser sandbox — bash tools (e.g., curl or other CLI operations) are not available.

Q: Can I install JDK in the conda environment?

Yes. Your code runs in an independent conda environment (with customizable dependencies), and browser interaction connects to the organizer-provided cloud sandbox via CDP. We recommend verifying that your custom environment runs stably under conda on your own machine first.

Q: Are there restrictions on programming languages or open-source frameworks?

No restrictions, but you must standardize your input/output interfaces according to the competition template requirements, and your code must run stably in a conda environment.

Q: Where are downloaded files stored in the sandbox? What's the temp path? Disk space limit?

Use Playwright-supported download methods to save files to your code directory — the path is defined by you. There's no disk space limit for download operations.

Network & Security

Q: Will websites fail to load due to geo-blocking during the competition?

The evaluation sandbox is provided by the organizers and will try to avoid geo-blocking/regional restriction issues.

Q: My API service only has an IP, no domain. What about domain whitelisting?

If your agent calls the OpenAI official API, the organizers have already set up acceleration for overseas access. You can configure `config.json` to call OpenAI's API directly without purchasing a domain.

Please provide the actual domain your Agent calls at runtime for whitelisting. Submit here: <https://ml.jsjform.com/f/jXbK5J>

Q: Do I need to put my API Key in plaintext in the Git repo? Is my code safe in the repo?

Your private repo is only visible to your team. The evaluation environment has strict security policies to protect your API Key.

After the competition, the organizers will either destroy the private repo or transfer it to you, based on your preference.

3. Competition Rules & Constraints

Steps & Execution

Q: Is there a time limit per question?

There's no per-question time limit during evaluation. Teams should plan according to the overall evaluation time limit.

Q: What exactly is the max steps limit?

The official maximum is 100 steps. Teams can set their own max step count based on overall time limits and Agent strategy, as long as it doesn't exceed 100.

Q: How is "one step" defined? Is screenshot + thinking + action = 1 step? Or each action call = 1 step?

A Web Agent predicting an action based on current page content and executing it counts as one step.

Q: After exceeding max steps, is it scored as failure or evaluated at current state?

Exceeding 100 steps results in immediate failure.

Q: Are there timeout limits per model call or call count limits?

Each model call must not exceed 3 minutes. There's no limit on call count. Your model must support 8 concurrent process requests.

Q: Is there an overall evaluation timeout?

Total evaluation time must not exceed 4 hours.

Operation Boundaries

Q: Can I navigate to a search engine to search?

No. External search engines are not allowed. You may use the website's internal search.

Q: Is direct CDP manipulation allowed?

Direct CDP manipulation is not allowed. You must create a CDP session based on Playwright, then execute CDP operations from there.

Q: For tasks involving PDFs, can I download and process them with a coding agent?

Yes. After navigating to the page, downloading the data and analyzing it to get the answer counts as one step overall.

Q: Will tasks require login to complete?

No. Competition questions will try to avoid requiring login. If login requirements appear in the open-source dataset, it may be due to website changes.

Q: Will CAPTCHAs appear during the competition?

Anti-bot/risk-control verification (e.g., click verification): Possible — this is part of the real web environment and your Agent needs to handle it. The organizers will try to select websites with low risk-control levels.

Models & Compliance

Q: Are there restrictions on which models teams can use?

Agents must call models via OpenAI-compatible API format. Two options:

Option A: Commercial Closed-Source API

Maximum allowed versions per provider:

- OpenAI — GPT-5.4
- Anthropic — Claude 4.6
- Google — Gemini 3.1
- xAI — Grok 4.3
- Zhipu AI — GLM-5V-Turbo
- Moonshot — Kimi-K2.6
- Alibaba Cloud — Qwen3.7 (new)

Versions beyond those listed above are prohibited. For providers not listed, only models released before July 16, 2026 are allowed (based on official public release announcement date). Models released after that date are prohibited.

Option B: Self-developed / Self-deployed Models

Allowed. Model weights must be submitted to the organizers for verification after the competition.

The organizers will audit model usage and strictly investigate any wrapper services or relay bypasses of version restrictions. Violations will result in disqualification.

Q: Can I mix multiple models (e.g., GPT for reasoning, Claude for OCR)?

Yes, but all models must comply with the organizer's version restrictions.

Q: If a model API call hits 429 rate limit or times out, does that step fail or can I retry?

You can retry — this needs to be implemented in your Agent design.

Q: Can a team submit multiple Agent versions for A/B testing?

There's no limit on code submission versions or frequency, but formal evaluation is limited to 1 attempt per day.

4. Scoring Criteria

Scoring & Ranking Logic

Q: What's the difference between automated evaluation scores and final published scores?

Automated evaluation score: The score provided by the bot during the competition is only the task answer accuracy score (whether answers are correct), and does not include trajectory scoring.

Final published score: After the evaluation window closes, final scores combining trajectory correctness will be published after full review.

Q: What are the final scoring criteria?

The competition uses Protocol III (navigation + data extraction) scoring:

1. Whether the agent navigated to the target page
2. Whether the correct field was extracted (semantic equivalence is sufficient)

Final results will be reviewed by human judges.

Q: What determines the final leaderboard ranking? Highest score or most recent?

Final ranking is sorted by composite score (tasks where both answer and trajectory pass). Higher composite score ranks first; at the same score, fewer total steps ranks first; at the same score and steps, less total time ranks first.

If a team has multiple submissions, the highest composite score is used for ranking.

Q: Will all teams' scores be reviewed?

Yes. The organizers will review every team's results, combining trajectory + answer scoring, and provide each team with their highest score after the review period.

Q: Does judging use the participant's model or a unified model?

Participants only need to build their Agent. The scoring pipeline is run by the organizers — you don't need to configure a judging model.

Answer Requirements

Q: What's the format requirement for agent_answer? Plain text? JSON? How to handle multiple fields?

Plain text. Multiple fields can be separated by \n or spaces.

Q: How strict is "semantic equivalence" — must numbers match exactly or is rounding OK? Are CN/EN or simplified/traditional equivalent?

Based on the question's requirements and the values shown on the target page. Languages are equivalent — there's no language restriction.

Q: Is partial completion scored? (Correct navigation but wrong answer, vs wrong navigation but correct answer)

Only when both navigation and data extraction are correct is the task considered complete. All other cases count as failure.

5. Troubleshooting

Smoke Test Failures

#	Cause	Solution
1	Cannot connect to sandbox after receiving CDP_URL	Keep the browser connection compatibility logic from the competition template repo
2	capture.json not properly capturing network requests for some tasks	Check your network request capture logic
3	Cleared all tabs after connecting to sandbox, causing browser to close	Don't close all tabs
4	Model call failed	Check API Key configuration and network connectivity
5	No response during execution causing timeout	Check if your model call has a timeout retry mechanism

Local vs. Sandbox Differences

Q: My code works locally with Playwright/Puppeteer but behaves differently in the sandbox. What do I do?

廖雨亭(Leo) 0135
2026-8-21

廖雨亭(Leo) 0135
2026-8-21

The sandbox connection has slight differences from local connections. The competition template repo's browser initialization and connection sections have been made compatible — make sure to keep that logic.

廖雨亭(Leo) 0135
2026-8-21

廖雨亭(Leo)
2026-8-

廖雨亭(Leo) 0135
2026-8-21

廖雨亭(Leo) 0135
2026-8-21

廖雨亭(Leo)
2026-8-21

廖雨亭(Leo) 0135
2026-8-21

廖雨亭(Leo) 0135
2026-8-21

廖雨亭(Leo)
2026-8-21

廖雨亭(Leo) 0135
2026-8-21

廖雨亭(Leo) 0135
2026-8-21

廖雨亭(Leo)
2026-8-21